

Introduction to Machine Learning

CS4375 --- Fall 2014

Course Project

1

Overview

- 15% of course grade
- spans the rest of the semester: you should start **now**
- provides hands-on experience with employing machine learning techniques to tackle a prediction task
 - However, you may choose not to use machine learning
- everyone will work on the same task
 - you may work in a team of up to three people

2

Task: Binary Prediction Task

- Class value: 0/1
- 4 discrete-valued features
 - 3 binary-valued, 1 six-valued
- 172 real-valued features
- Training set: 12,000 instances
- Test set for preliminary evaluation: 5,000 instances
- Test set for final evaluation: 8,000 instances
- The test set will not contain targets, of course

What you need to do

- Predict the class values of the test instances
 - You may use any algorithms and techniques you want, including those that are not introduced in this course
 - You may even employ non-learning-based methods
 - You may use any publicly available software packages

Key Dates

- Preliminary evaluation: Sunday, Nov 23
- Final evaluation: Sunday, Dec 14
- Project report: Wednesday, Dec 17

Preliminary Evaluation (Due: Nov 23)

- 4% of the course grade
- The test set for the preliminary evaluation will be available on Nov 16 (one week before the deadline)
- Need to submit (1) your prediction file, and (2) a README file listing the names of all group members to eLearning

Prediction File Format

- The file should be plain text
- Each line should contain **only** the predicted class value (0 or 1) of a test instance
- # lines in prediction file should be the same as # lines in test file
- You must return predictions to us in the same order as the instances in the test file
- Sample prediction file:
0
0
1
...

Scoring

- We will compute the accuracy (i.e., percentage of correctly classified test instances) of the system submitted by each team
- We will rank the systems based on their accuracies
- Your score in the preliminary evaluation will be based on your rank
 - $\text{Score}(s) = (\text{total \# of teams} - \text{rank}(s) + 1) / (\text{total \# of teams})$
 - So the scores will be evenly distributed between 0 and 1

Final Evaluation (Due: Dec 14)

- 7% of the course grade
- **Goal:** a second chance for you to improve your system after seeing how your team performs relative to the other teams
- The test set for the preliminary evaluation will be available on Dec 7 (one week before the deadline)
- Need to submit (1) your prediction file; (2) your code; and (3) a README file containing instructions on how to compile and run your system, as well as a listing the names of all group members to eLearning
 - Prediction file format and scoring are the same as those in the preliminary evaluation

Project Report (Due: Dec 17)

- 4% of the course grade
- The project report should describe everything that the team did for the project. It should
 - include approaches that were attempted but were ultimately not employed because of their poor performance, for instance
 - the approach that was chosen for final evaluation
 - lessons learned
- More details on the project report later ...

Training Set

- Available from the course website

Challenges

- Parameter tuning
 - Almost all learning algorithms have their own set of parameters
 - E.g., for neural nets, we need to specify the number of hidden layers and the number of hidden units per layer
 - The performance of a learner is to a large extent determined by these parameters
 - You probably want to tune them on validation data

What else can you try?

- Different learning algorithms
 - You can employ multiple learning algorithms to make predictions
 - Some learners may perform better than others
 - You may consider discarding the bad ones or putting less weights on them
 - How can I determine which learners are bad?
 - Use your validation data
 - How do I know how much weight I should give to a learner's prediction?
 - Use your validation data

Anything else?

- Feature selection
 - Motivation: using all the available features may not always yield better results than using a subset of them
 - While many learners can (implicitly) select features, in many cases it may be good to explicitly identify and filter out the irrelevant features
 - How can we select good features?
 - Information gain
 - Let the decision tree learner tell you
 - Other methods (consult the literature)

Summary

- You can use whatever approach you want to do make predictions
 - Be creative!
- You can even discuss your approach with other groups
 - But you are not allowed to share your team's predictions with other groups